

AI Visibility Audit

A Product Case Study on Measuring and Improving Brand
Visibility in AI Search

Suchithra R

Product Management Portfolio / September 2026

A structured audit of how three AI assistants respond to five buyer-intent queries for a specialty coffee brand.

02 / EXECUTIVE SUMMARY

The product in one minute

A repeatable way to inspect brand presence across AI answers, then turn gaps into a sequenced action plan.

THE CORE QUESTION

When a buyer asks an AI assistant for relevant options, does the brand appear, get described accurately, and have credible supporting sources?

AT A GLANCE / FACTS SHOWN ON THE SITE

5

buyer queries

Across five intent stages

3

AI engines

ChatGPT, Gemini, Perplexity

6

dimensions

Defined in the framework

5

fixes

Shown as expandable cards

WHAT THE READER GETS

Evidence of visibility

A 5 x 3 scorecard captures mention, placement, and citation quality at the query level.

A decision path

The page converts gaps into five recommendations grouped by impact, effort, and sequence.

Scope note: the live experience is a published audit case study. It does not accept a new brand or run an audit for visitors.

03 / THE PROBLEM

Discovery is harder to see

A conventional search position does not reveal whether an AI answer will recommend, omit, or misdescribe a brand.

Problem

Teams need to know whether AI assistants surface the brand for buyer questions that matter, not only for questions that already contain its name.

User pain

Without a query-level view, a team can see a strong branded answer yet miss a category discovery gap or a weak citation behind a favorable answer.

Product opportunity

Make AI visibility observable by intent, model, and evidence source, then connect each gap to an action the team can prioritize.

WHY EXISTING SEARCH METRICS ARE INCOMPLETE

Rank and traffic describe search result pages and site visits. This audit examines the answer itself: whether the brand is named, how it is positioned, and which sources support the response. These are complementary views of discovery.

Case study question: How can a company understand whether AI systems surface, recognize, and represent its brand when buyers ask relevant questions?

04 / USER + HYPOTHESIS

Who needs this view?

The site presents a client brief, content changes, structured data work, and an execution roadmap. These suggest the roles below.

Product / growth marketing

Needs to know which buyer intents do not surface the brand. Can decide where to focus category messaging and campaign content.

SEO / content teams

Needs to see how pages, descriptions, and sources relate to AI answers. Can prioritize content and structured information updates.

Brand or marketing lead

Needs to understand whether answers match intended positioning and whether favorable mentions rest on credible sources. Can allocate on-site versus external authority work.

PRODUCT HYPOTHESIS / PROPOSAL

If teams review the same buyer-intent questions across AI assistants and record mentions, descriptions, and sources, they can find visibility gaps and choose better-targeted actions.

INPUT

Five buyer-intent questions

ANALYSIS

Cross-engine comparison

OUTPUT

Scorecard + findings

DECISION

Which content, source, or positioning gap to address first

User roles are inferred from the actions shown on the site; no user interviews or personas are documented.

05 / PRODUCT WORKFLOW

From question to action

The site describes a seven-step manual audit. The visitor can inspect its results and expand the five recommendation cards.

01

Frame buyer intent

Choose discovery, comparison, consideration, use-case, and trust questions.

02

Run identical prompts

Ask ChatGPT, Gemini, and Perplexity the same five questions.

03

Collect responses

Record response text, mentions, positioning, and cited sources.

04

Check accuracy

Compare claims and descriptions with the brand's offering.

05

Find visibility gaps

Locate absences, misrepresentation, and competitor displacement.

06

Prioritize fixes

Order gaps by stated impact, effort, and dependency.

07

Deliver recommendations

Present content, schema, and external-source actions.

Interaction boundary: the page is a read-only case study; the recommendation cards expand to reveal problem, solution, and expected impact.

06 / AUDIT METHODOLOGY

A small, structured sample

The published method uses five distinct buyer questions across three AI assistants, producing a 15-cell comparison grid.

TEST DESIGN / FACT

Discovery	Best South Indian filter coffee brand online
Comparison	Brand vs competitor for daily drip coffee
Consideration	Best subscription for trying different roasts
Use-case	Is the brand good for cold brew?
Trust	Why is the brand considered good?

CAPTURE + INTERPRET

The site says responses and cited sources were recorded in a structured database, then assessed for mention, prominence, accuracy, positioning, and category fit. It compares the three answers per intent and maps gaps to fixes.

WHAT IS NOT DOCUMENTED

The page does not provide the underlying response logs, timestamps for individual runs, sample repetitions, a calibrated scoring rubric, or a numeric formula. Its outcome labels should be read as reported qualitative judgments from this sample.

Model names are listed on the site as GPT-4o, Gemini 1.5 Pro, and Perplexity AI. This report does not re-run the historical audit.

Measure the answer, not just a rank

The site names six dimensions. It does not publish dimension values or an overall score calculation.

DIMENSION	WHAT IT MEASURES	WHY IT MATTERS / INTERPRETATION
Visibility	Mention or absence	A brand must appear before it can be chosen.
Accuracy	Correct, current description	Wrong claims can weaken buyer trust.
Citation quality	Source relevance and strength	A favorable answer may rest on weak evidence.
Intent matching	Fit to buyer's question	A mention is less useful if it misses the task.
Brand positioning	How the brand is framed	The answer may differ from intended positioning.
Authority signals	Structured and external evidence	Signals may help support category association.

WHAT CAN BE SCORED TODAY?

Reported query-level outcomes

For example: absent in 3/3 discovery answers; top or second on subscription.

No validated 0-100 score

An overall score appears only as a future SaaS idea, with no published formula.

Treat the named dimensions as a framework and the 5 x 3 scorecard as a qualitative snapshot, not a scientific index.

08 / AUDIT RESULTS

The reported 5 x 3 scorecard

The image below is a cropped capture of the live scorecard. The simplified reading beside it uses the site's labels.

QUERY	INTENT	CHATGPT	GEMINI	PERPLEXITY
"Best South Indian filter coffee brand online"	Discovery	⊗ Absent	⊗ Absent	⊗ Absent
"Blue Tokai vs Third Wave – daily drip coffee"	Comparison	✔ Wins	✔ Wins	⚠ Wins*
"Best coffee subscription for trying different roasts"	Consideration	✔ #1 pick	⚠ 2nd place	⊗ Absent
"Is Blue Tokai good for cold brew?"	Use-case	✔ Strong	✔ Strong	⚠ Correct*
"Why is Blue Tokai considered a good brand?"	Trust	✔ Strong	✔ Strong	⚠ Correct*

SOURCE CAPTURE

The site groups each question by buyer intent and compares ChatGPT, Gemini, and Perplexity. Green denotes stronger mentions; amber marks weaker cited sources; red marks absence.

0 / 3

Generic discovery

Brand absent in all three answers

3 / 3

Branded comparison

Reported as a win in all three

1 / 3

Subscription leader

First in ChatGPT only

OTHER REPORTED PATTERNS

01

Use-case + trust

ChatGPT and Gemini are marked strong; Perplexity is marked correct with weaker source quality.

Source: published scorecard and findings sections. These are site-reported observations, not independently reproduced measurements.

09 / KEY INSIGHTS

What the pattern suggests

Each interpretation below follows from the reported answer patterns. Causal explanations remain hypotheses.

01 / DISCOVERY

Category intent is the gap

Observation: the brand is absent from all three generic filter-coffee answers, while winning the branded comparison.

Implication: audit branded recall and non-branded discovery separately; prioritize category relevance signals.

02 / CONSIDERATION

Visibility changes by intent and engine

Observation: subscription is first in ChatGPT, second in Gemini, and absent in Perplexity.

Implication: test subscription messaging against variety-seeker questions before assuming broad visibility.

03 / EVIDENCE

A mention can still be fragile

Observation: the site flags Perplexity's supporting sources as weaker on several favorable answers.

Implication: track where the answer's evidence comes from, not only whether the brand appears.

The site calls missing schema a likely root cause. The audit does not isolate schema's effect, so this report treats that link as a testable hypothesis.

10 / RECOMMENDATIONS

Sequence the fixes

The five existing actions are ordered here by the site's impact and effort labels, with dependencies made explicit.

1

Clarify category content**IMPACT HIGH / EFFORT LOW**

Gap: Discovery absence. A dedicated filter-coffee page and product language can state category fit clearly.

2

Add structured information**IMPACT HIGH / EFFORT LOW**

Gap: Missing on-site signals. Product and organization markup may help machines interpret offerings; validate with re-audit.

3

Reframe subscription copy**IMPACT HIGH / EFFORT MEDIUM**

Gap: Inconsistent consideration. Make roast variety easier to understand for variety-seeker queries.

4

Build credible external coverage**IMPACT HIGH / EFFORT HIGH**

Gap: Weak cited sources. Expert reviews and relevant roundups could broaden supporting evidence.

5

Expand category presence**IMPACT MEDIUM / EFFORT HIGH**

Gap: Sustained discoverability. Category-specific education and listings can follow the on-site foundation.

Impact labels are product judgments shown on the site, not measured lift. The site's time-to-visibility forecasts are omitted because no validation evidence is provided.

11 / PRODUCT OPPORTUNITIES

From report to repeatable product

The current experience communicates one audit. A future product could make that evaluation repeatable and comparable.

CURRENT PRODUCT / FACT

A public case-study page with a manual seven-step method, a five-query scorecard, six evaluation dimensions, findings, a priority matrix, and five expandable recommendation cards.

Visitor interaction is limited to navigation and expanding cards; there is no live audit intake, account, dashboard, or historical monitoring workflow shown.

POTENTIAL FUTURE PRODUCT / PROPOSAL**1. Repeatable runs**

Save prompts, models, run dates, and raw answers to compare like with like.

2. Query-level trends

Show whether a visibility gap persists across repeated runs.

3. Competitor context

Compare brands within the same buyer-intent question set.

4. Citation review

Group cited sources by type, relevance, and ownership.

5. Action tracking

Connect a recommendation to a re-audit and observed movement.

Suggested order follows a measurement dependency: record the evidence first, then show trends and evaluate whether actions coincide with movement.

How to measure product value

These are proposed product metrics for a launched audit tool. They are not measured by the current page.

PROPOSED NORTH STAR

Teams completing a repeat audit and making at least one documented decision from it

Acquisition	Qualified audit starts Shows whether the problem attracts relevant teams.
Activation	Completed first audit Confirms that a team reaches a usable baseline.
Engagement	Findings reviewed; actions saved Shows whether output informs work.
Retention	Repeat audits in a defined interval Tests whether monitoring is worth returning for.
Outcome	Gap closure by intent over time Checks whether the tracked problem changes.

Guardrail: report score movement with the underlying prompts, model versions, and source evidence so a changed score can be interpreted.

13 / LIMITATIONS + CAVEATS

Treat a snapshot as a signal

The audit is useful for finding questions to investigate. It is not proof of a stable market position or a causal fix.

01

Small query set

Five prompts cover selected intents, not the full range of buyer language.

02

Response variability

AI answers can differ across runs, accounts, time, and model updates.

03

Prompt sensitivity

Minor wording changes can change which brands appear and in what order.

04

Source transparency

Raw transcripts and run timestamps are not available on the public page.

05

Unspecified scoring

Six dimensions are named, but no numerical rubric or formula is published.

06

Causality

A missing category mention does not by itself prove that schema, content, or citations caused it.

Next validation step: repeat a frozen prompt set, archive raw answers, then compare changes after each action.

14 / PRODUCT THINKING TAKEAWAYS

Turn ambiguity into decisions

The strongest part of this case is the path from an unclear discovery problem to a compact, discussable evidence set.

01

Decompose the problem

Separate branded recall, category discovery, intent fit, and source quality.

02

Design a repeatable lens

Use the same buyer questions across assistants so differences are legible.

03

Keep evidence close

Show the query-level answer pattern before interpreting it.

04

Prioritize action

Sequence content, structured information, and external evidence work.

05

Close the learning loop

Re-run the audit to test whether the proposed fixes coincide with change.

EVIDENCE KEY

FACT = shown or stated on the live page. INSIGHT = an interpretation of its reported answers. PROPOSAL = a future action, product feature, or measurement plan.

Source reviewed: live AI Visibility Audit case-study page, including the scorecard, method, findings, priority matrix, and expanded recommendation cards; 26 September 2026. No raw audit log was available on the page.