


Myntra Personal Shopping Assistant

AI-powered recommendations, price tracking and smarter wishlists to reduce decision fatigue and drive repeat purchases



Suchithra R

Product Management portfolio · September 2026

Independent case study. Not affiliated with or endorsed by Myntra. Targets, baselines and sizing are illustrative assumptions.



Contents

Executive summary 3

1. Objective and ideal future 3

2. Problem and opportunity 4

3. Target audience 4

4. Solution and scope 5

5. Core flows 6

6. Success metrics 8

7. Prioritization and phasing 10

8. Experiment plan 10

9. Risks and mitigations 11

10. Design guidelines 11

11. Assumptions and open questions 11

Executive summary

One assistant that helps shoppers decide, waits for the right price, and turns the wishlist into a plan.

+25%

repeat customers within 3 months

+20%

session time

+15%

wishlist adds per user

+10%

revenue from personalized picks

- **The problem.** Shoppers face too many options, miss price drops because nothing tracks them, and use a wishlist that only stores items. Purchases get delayed or abandoned.
- **The solution.** An AI shopping assistant with three surfaces: a **“Picked for you”** carousel with reasons, **price-drop alerts** on wishlist items, and a **smart wishlist** with categories and “complete the look” suggestions.
- **The approach.** Earn trust with utility first (alerts and wishlist organization), then layer on personalization. Ship in three phases, prioritized with RICE and proven in A/B tests with holdouts.
- **Who it is for.** Frequent shoppers, fashion enthusiasts, price-conscious buyers and wishlist creators. Launch focus is on wishlist creators and price-conscious buyers, where the lift is most direct.
- **The biggest risk.** Price alerts can teach users to wait for discounts and erode margin. The plan limits alerts to meaningful drops, caps frequency, and tracks gross margin per order as a guardrail.

1 Objective and ideal future

What we are building and what success looks like.

Objective. Create an AI-driven personal shopping assistant that increases engagement, drives repeat purchases and builds loyalty by offering personalized product recommendations, price tracking and wishlist management.

The ideal future

- Myntra becomes the go-to destination where personalized recommendations show up effortlessly.
- Decision fatigue drops, and timely price-drop alerts convert into more purchases.
- Loyalty grows because Myntra becomes indispensable for managing everyday shopping needs.

One-line vision. Make every visit to Myntra easier than the last: fewer choices, better timing, and a wishlist that works for the shopper.

How the assistant delivers each outcome

Desired outcome	Mechanism	Where the user meets it
Effortless personalization	Ranked picks from behaviour and preferences, each with a short reason	Homepage and product-page carousel
Less decision fatigue	A short list of relevant items instead of an endless grid	Carousel, category view of the wishlist
More purchases from timely alerts	Price tracking on saved items with user-controlled thresholds	Bottom sheet, push and email
Loyalty and indispensability	Wishlist becomes a planning space with categories and look-building	Wishlist tab

2 Problem and opportunity

Three current problems, the likely root causes, and the opportunity each one opens.

Problem	What the shopper experiences	Likely root cause	Opportunity
Overwhelming choices	Endless grids, too many similar options, hard to decide	Relevance is diluted; little explanation of why an item is shown	Personalization with visible reasons reduces decision fatigue
Missed opportunities	Price drops happen unnoticed; the item sells out or the deal ends	No tracking or alert mechanism for saved items	Timely alerts turn passive wishlists into active purchases
Unintelligent wishlist	A long, unsorted list that is hard to act on	Wishlist stores items but does not organize or suggest	Categories and look-completion make it a planning tool

Problem statement. Frequent shoppers on Myntra face too many choices, cannot act on price changes, and have a wishlist that only stores items. The result is delayed or abandoned purchases and weaker loyalty.

Hypotheses to validate

#	If we...	Then...	Because...	Primary metric
H1	Show a short list of picks with reasons	Shoppers add to wishlist or bag more often	Relevance plus explanation lowers effort	Wishlist adds per user
H2	Offer price-drop alerts at the moment of saving	More saved items become purchases	Timing removes the need to keep checking	Wishlist-to-order rate
H3	Let users organize and get look suggestions	Basket size and return visits rise	The wishlist becomes a reason to come back	Repeat customers in 3 months

Tension to manage. Less decision fatigue should shorten time to decision, yet the target asks for longer sessions. We treat session time as an engagement signal, and track time to first add-to-bag as a guardrail.

3 Target audience

Four segments, what each needs, and how the assistant serves them.

Segment	Who they are	Core need	What the assistant does	Signal we track
Frequent online shoppers	Regular users who want a streamlined experience	Speed and relevance	Carousel and faster paths to items they like	Sessions per month, add-to-bag rate
Fashion enthusiasts	Trend followers who build curated looks	Inspiration and coordination	Complete-the-look suggestions	Items per order, look adds
Price-conscious buyers	Compare prices and wait for discounts	Confidence in the price	Price-drop alerts with old and new price	Alert opt-in, alert-to-purchase
Wishlist creators	Save many items but delay buying	A nudge and a way to decide	Categories, alerts and suggestions	Wishlist-to-order rate

Launch focus. Wishlist creators and price-conscious buyers. They already show intent through saved items, so alerts and organization can lift conversion quickly.

4 Solution and scope

One assistant, three surfaces, built to feel like part of the app rather than an intrusion.



Scope


Scope	Item	Why	Phase
In scope	AI-based personalized product recommendations	Directly addresses decision fatigue and revenue from recommendations	2
In scope	Price tracking and notification system	Fastest path from saved item to purchase	1
In scope	Enhanced wishlist management (categories, AI suggestions)	Turns the wishlist into a retention surface	1 to 2
May be in scope	Social sharing of wishlists	Adds virality but is not needed to prove core value	3
May be in scope	AI-powered chatbot for style guidance	High effort and uncertain quality; revisit with data	3
Out of scope	Cross-platform shopping assistant	Needs external integrations and partnerships	n/a
Out of scope	Voice-based AI interactions	Not central to the current problem	n/a
Out of scope	Virtual try-on	Separate product investment	n/a

5 Core flows

Each flow shows the trigger, what the assistant does, and the low-fidelity screens that go with it.


Flow 1: Product recommendations

Trigger	The user lands on the homepage or a product page.
Assistant action	Shows a “Picked for you” carousel of 6 to 10 items ranked from browsing, search, wishlist and purchase signals, each with a one-line reason.
User action	Save to wishlist (heart) or view details.
Outcome	Fewer, more relevant choices. New saves feed the price-alert flow.
Guardrails	One compact module, not a takeover. Users can hide a pick (“Not for me”) to teach the model. New users see trending items in their size and category.




Flow 2: Price alerts

Trigger	The user saves a product to the wishlist, or taps the alert pill on an existing item.
Assistant action	Offers a choice: any drop, 10% or more, or a target price, and a channel (push, email or both). The tracker checks prices on a schedule.
User action	Turn on alerts, then tap Buy now from the notification.
Outcome	The wishlist becomes active. A drop sends a deep link back to the item with the user's size still selected.
Guardrails	Alerts show old and new price and stock status. Frequency caps apply per item and per week. No alerts for sizes that are out of stock.



Flow 3: Wishlist management

Trigger	The user opens the wishlist.
Assistant action	Suggests auto-categories, lets the user create their own (for example Office Wear, Casuals), and offers “Complete your look” inside a category.
User action	Add, remove and organize items; view AI suggestions; add the look to the bag.
Outcome	The wishlist becomes a planning space that raises basket size and repeat visits.
Guardrails	Suggestions appear when the user opens a category, not as pushes. Auto-categories are suggestions the user can ignore.



The wireframes are low fidelity and show structure and flow only. Visual design should follow the existing Myntra app language.


6 Success metrics

The four key results, how each is measured, and the guardrails that protect the business.

Goal	Metric	Target	How measured
Engagement	Session time	+20%	Average session duration for users exposed to the assistant versus a holdout, over 90 days
Engagement	Products saved per user	+15%	Wishlist adds per active user versus holdout
Repeat purchases	Repeat customers within 3 months	+25%	Share of buyers placing a second order within 90 days
Revenue	Revenue from personalized recommendations	+10%	Revenue from assistant-surfaced items (7-day attribution window)
Loyalty	NPS among frequent users	50+	In-app survey to users with 3 or more orders in 6 months (promoters rate 9 to 10)

Baselines are captured in the two weeks before launch. All targets are relative lifts measured against a holdout group.

Metrics tree



Guardrails

Guardrail	Why it matters	Acceptable range
Return rate	Better matching should not increase returns	No rise versus holdout
Notification opt-out rate	Alerts must stay welcome	No rise versus holdout
Gross margin per order	Alerts should not train users to wait for discounts	No drop versus holdout
Time to first add-to-bag	Confirms less decision fatigue, not more browsing	No rise versus holdout
“Not for me” rate on picks	Signals recommendation quality	Falls over time

7 Prioritization and phasing

RICE scores for every candidate, then a phased plan that builds trust before personalization.

Reach is the share of active users touched per quarter. Impact: 0.5 low, 1 medium, 2 high. Effort is in person-months. $RICE = Reach \times Impact \times Confidence \div Effort$. All inputs are illustrative assumptions.

Item	Reach	Impact	Confidence	Effort	RICE
Price-drop alerts	35	2	80%	2	28.0
Wishlist categories	35	1	90%	1.5	21.0
AI recommendations carousel	90	1	60%	4	13.5
Complete-the-look suggestions	25	2	50%	3	8.3
Social sharing of wishlists	10	1	50%	2	2.5
Style-guidance chatbot	15	2	30%	6	1.5

Why alerts and wishlist first

They score highest, are low effort, and reach users who already show intent. They also generate the signals (saves, categories, price sensitivity) that make the later recommendation and look-completion features better. The recommendations carousel scores lower on RICE than its strategic weight suggests, because it likely builds on existing ranking systems; that is an assumption to confirm with engineering.

Phased roadmap

Phase	Scope	Goal	Exit criteria
1 (weeks 0 to 8)	Price-drop alerts, wishlist categories, event tracking, baselines	Activate the wishlist	Wishlist-to-order lift; opt-outs and margin within guardrails
2 (weeks 8 to 16)	Picked-for-you carousel, complete-the-look suggestions	Reduce decision fatigue, grow basket size	Add-to-bag and revenue per session up; time to first add-to-bag not up
3 (weeks 16+)	Social sharing and style chatbot, only if Phase 1 to 2 succeed	Extend engagement	Go or no-go from experiment results

8 Experiment plan

Every feature ships behind a test with a holdout and pre-agreed guardrails.

Feature	Hypothesis	Design	Primary metric	Guardrails	Runtime
Price alerts	Alerts at save time convert more saved items	50/50 A/B among wishlist savers: alert prompt versus none	Wishlist-to-order rate (30 days)	Gross margin, opt-outs	4 weeks
Wishlist categories	Organization increases return visits	50/50 A/B among wishlist users	Return visits per month	Wishlist load time	3 weeks
Picked-for-you carousel	Ranked picks with reasons raise add-to-bag	50/50 for 3 weeks, then a 10% long-term holdout	Add-to-bag rate, revenue per session	Time to first add-to-bag, return rate	3 weeks + holdout
Complete the look	Pairings raise items per order	50/50 A/B on category views	Items per order	Return rate	4 weeks

- Randomize at the user level and keep users in one arm across sessions.
- One primary metric per test; secondary metrics are read as directional.
- Read the repeat-purchase target at 90 days, so keep a small holdout through the full window.
- Decide ship, iterate or stop from rules written before launch, not from mid-test peeks.

9 Risks and mitigations

Risk	Why it matters	Mitigation
Users learn to wait for discounts	Alerts can cut full-price sales and margin	Alert on meaningful drops, track margin per order, exclude items already at deep discount
Notification fatigue	Too many pushes lead to opt-outs and uninstalls	Per-item and weekly caps, digest option, clear opt-out control
Cold start and poor early picks	New users lack behaviour signal	Trending and size-aware defaults; let users hide picks to teach the model
Filter bubble	Over-narrow picks reduce discovery	Reserve slots for exploration; monitor category diversity
Price data accuracy	A wrong price destroys trust	Verify at send time; show stock and size availability
Privacy and consent	Personalization relies on behavioural data	Clear consent and controls; review current Indian data-protection requirements before launch
Feels intrusive	Assistant must extend the journey, not interrupt it	Compact modules, user-triggered suggestions, easy dismissal

10 Design guidelines

Focus on simplicity and seamless integration with the existing Myntra app design.

- **Feel like an extension.** Reuse existing components and patterns; the assistant should read as part of the journey, not an overlay.
- **Be transparent.** Label AI picks and show a short reason so recommendations feel earned, not random.
- **Keep the user in control.** Easy dismiss, hide and opt-out on every module and alert.
- **Respect attention.** Suggestions appear on user action wherever possible; pushes are capped.
- **Design for clarity.** Price alerts always show old price, new price and availability.
- **Make it accessible.** Sufficient contrast, readable type sizes and clear tap targets on every new component.

11 Assumptions and open questions

Assumption	How to confirm	Owner
Current baselines for session time, wishlist adds and 90-day repeat rate	Analytics export in the two weeks before launch	Data / Growth
Recommendation ranking infrastructure can be reused	Engineering feasibility review	Engineering
Price history is available at product and size level	Data audit of catalogue and pricing feeds	Data / Catalogue
Push permission rates are high enough for alerts to matter	Check opt-in rates by cohort; add in-app inbox if low	Growth
Reach and effort estimates in the RICE table	Sizing with product analytics and engineering	Product